

# S.H.I.E.L.D

# Scudo Heap Implementation Exploits, Leaks, and Defenses

---

Tom Mansion



OFF-BY-ONE  
2025

Quarkslab

# Who I am

## Tom Mansion

- R&D Engineer @Quarkslab
- CTF Enjoyer
- Heap exploitation enthusiast
- Focused on Android




# Presentation outline

- 1. Introduction to heap allocators**
  - a. GLIBC's ptmalloc2 refresher
  - b. How can we exploit it?
- 2. Introduction to Scudo**
  - a. Scudo security principles
  - b. Scudo building blocks
- 3. Scudo exploitation techniques**
  - a. Use-after-frees and the quarantine
  - b. Zygote to the rescue!
  - c. Forged Commitbase
  - d. Safe Unlink
- 4. Conclusion**

# Introduction - what even is a heap allocator?


- Dynamic memory management runtime
- Lives in libc
- Used by almost every program
- Called very often



# Introduction - what even is a heap allocator?

Properties of a good allocator:


- Transparent



# Introduction - what even is a heap allocator?

Properties of a good allocator:


- Transparent
- Reliable



# Introduction - what even is a heap allocator?

Properties of a good allocator:

- Transparent
- Reliable
- Efficient








# Introduction - what even is a heap allocator?

Properties of a good allocator:

- Transparent
- Reliable
- Efficient
- **Fast**
- (Secure)



# GLIBC's ptmalloc2 (refresher)

## *Properties of ptmalloc2*

lower  
addresses

Space reserved  
for the allocator


← top chunk

wilderness

higher  
addresses

# GLIBC's ptmalloc2 (refresher)

## *Properties of ptmalloc2*




# GLIBC's ptmalloc2 (refresher)

## *Properties of ptmalloc2*

- Chunks of different sizes can be adjacent
- Deterministic

lower  
addresses




top chunk

higher  
addresses

# GLIBC's ptmalloc2 (refresher)

## *Properties of ptmalloc2*

- Chunks of different sizes can be adjacent
- Deterministic
- Free chunks can be stored in freelists
- Freelist metadata is stored on the heap



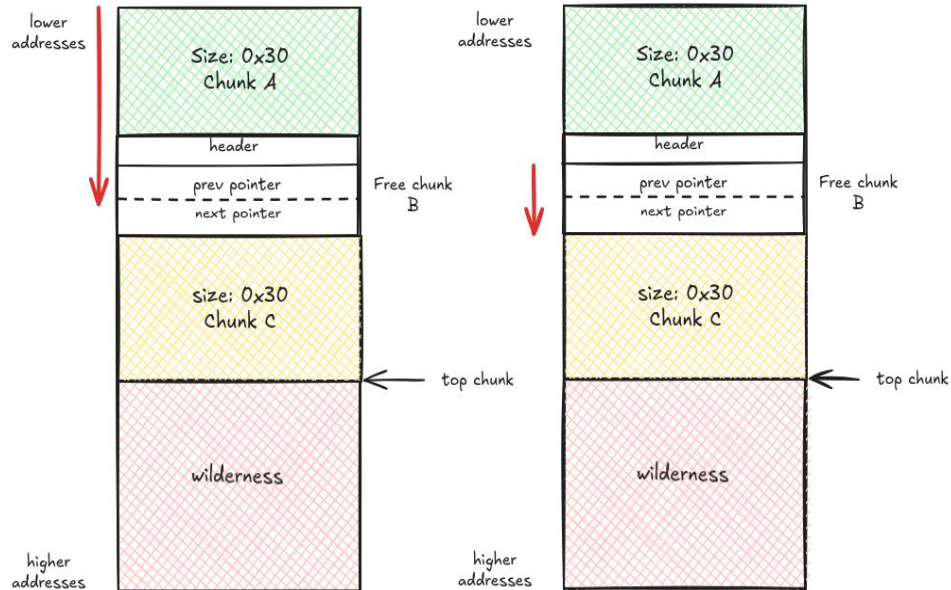
# GLIBC's ptmalloc2 - how can we exploit it?

Different classes of vulnerabilities (not exhaustive)

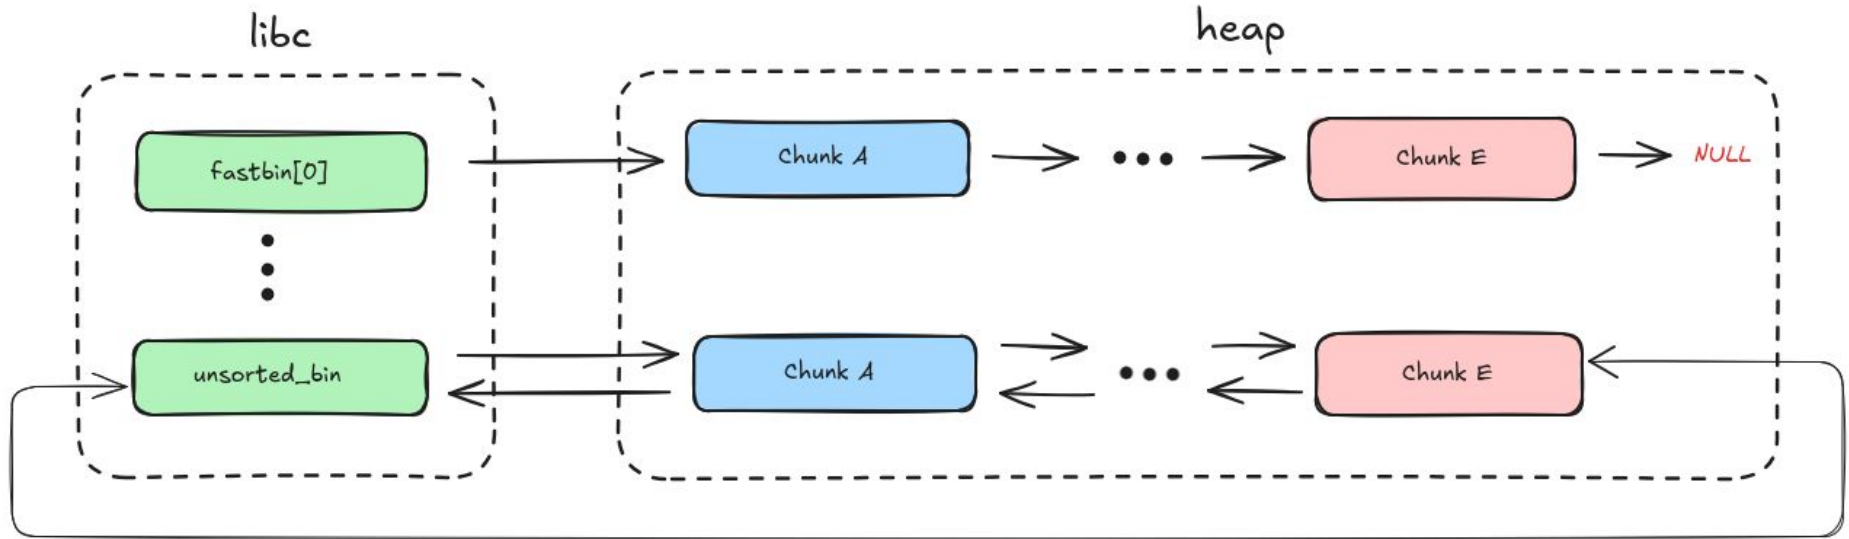
- Buffer overflow
- Use-After-Free
- Double-Free

**What can we gain control over?**

- User-controlled data
- Allocator control metadata




# GLIBC's ptmalloc2 - what metadata do chunks hold?



# Introduction to Scudo

- Default allocator on Android
- **Hardened**
- Configurable
- Part of the LLVM project





# Introduction to Scudo - Security principles

- **Separating allocation classes**
- **Introducing randomness wherever possible**
  - Randomizing the order of allocated chunks
- **Protecting the chunk headers with a checksum**
  - Prevents small buffer overflows
  - Based on a global secret stored in libc
- **Removing control metadata from the heap**
  - Prevents corruption of freelist pointers

# Scudo building blocks - Primary and Secondary allocators

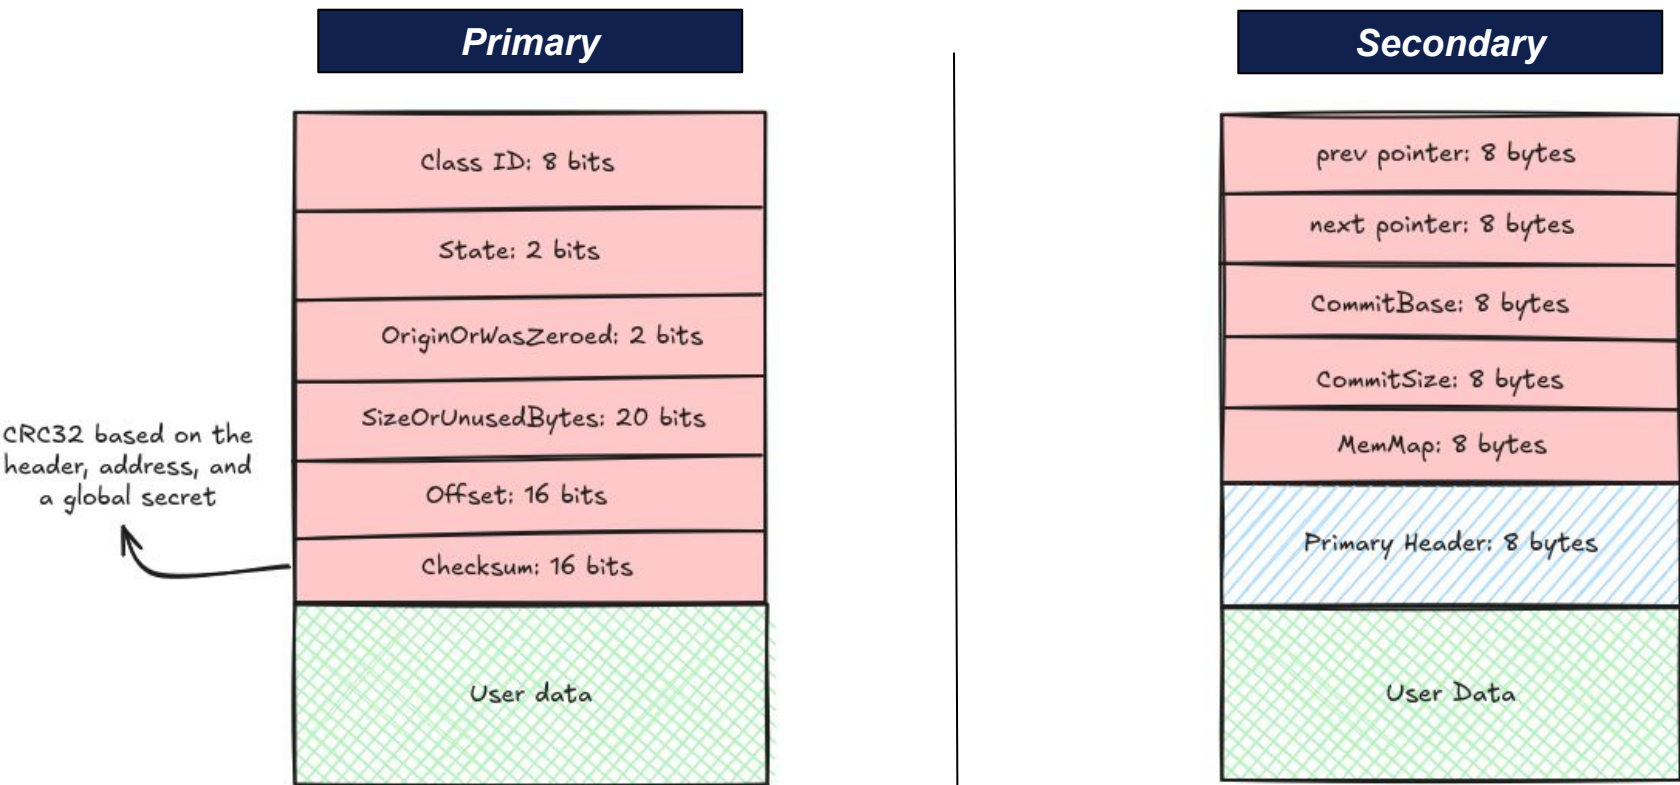
## *Primary*

- **Small** allocations (< 64kB)
- One Region per chunk size
- Local cache used for **everything**
- Creates chunks with a primary header
- (Has a quarantine cache)

## *Secondary*

- **Large** allocations (> 64kB)
- Creates chunks with mmap()
- Local cache for free chunks
- Creates chunks with a primary+secondary header
- (Has a quarantine cache)

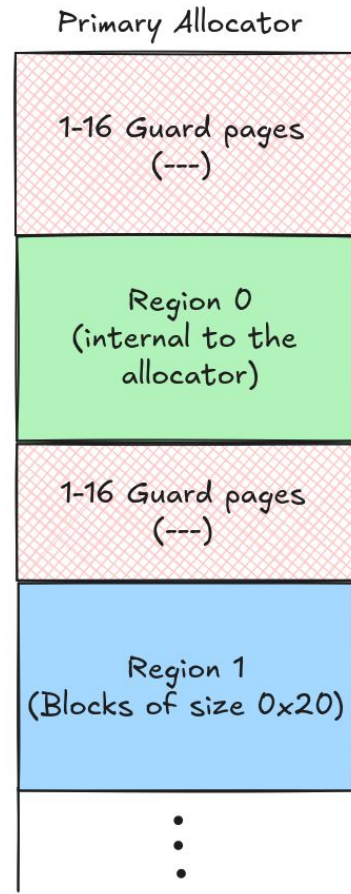
# Scudo building blocks - Chunk Headers



# Scudo building blocks - Regions & Guard Pages

*Scudo separates chunks by their size*

- Mmaps a bunch of pages for the Primary
- Divides it into 32 “Regions”
- Random amount of 1-16 guard pages for each Region



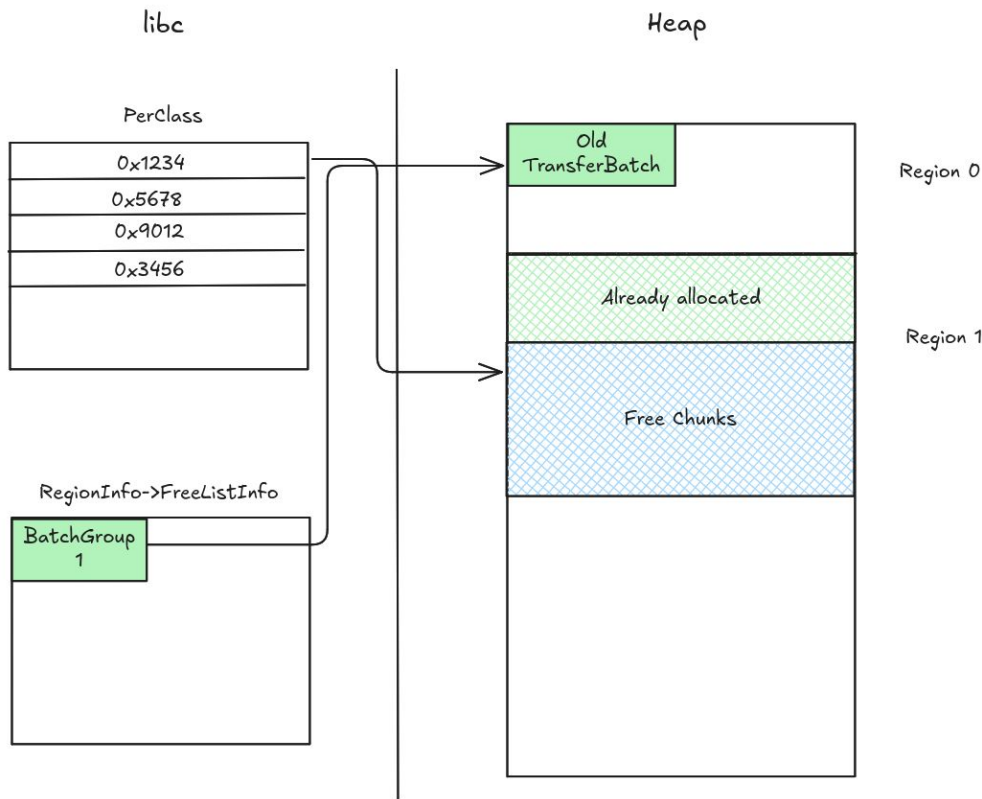
# Scudo building blocks - Local cache

- Free list implementation
- One per Region (PerClass)
- Fully stored in libc
- Drain/Fill operations use TransferBatches

```
struct alignas(SCUDO_CACHE_LINE_SIZE) PerClass {  
    u16 Count;  
    u16 MaxCount;  
    // Note: ClassSize is zero for the transfer batch.  
    uptr ClassSize;  
    CompactPtrT Chunks[2 * SizeClassMap::MaxNumCachedHint];  
};
```


# Scudo building blocks - Transfer Batches

- Part of a region's second freelist
- Cache layer between Region and local cache



# Scudo building blocks - ARM MTE support

- Scudo developed to support it
- Different tag for header and contents
- Different tag for Primary/Secondary header
- Memory is re-tagged upon deallocation



Source:  
<https://developer.arm.com/documentation/10803/5/0100/How-does-MTE-work-?lang=en>

*And now, the moment you've all been waiting for*




# The case of Use-After-Frees

- Without the Quarantine and MTE, Scudo doesn't mitigate UAFs!
- No inline metadata to corrupt
- Type confusion exploits are the new target
- Example from a basic CTF challenge

```
struct Dog {  
    char name[12];  
    void (*bark)();  
    void (*bringBackTheFlag)();  
    void (*death)(struct Dog*);  
};  
  
struct DogHouse{  
    char address[16];  
    char name[8];  
};
```

# Scudo Exploitation

- Forged Commitbase
- Safe Unlink
- Zygote to the rescue!
- Patches 🥲



# Forged Commitbase & Safe Unlink

Techniques published by P. Mao  
et al, in Usenix WOOT 24'

[Exploiting Android's Hardened  
Memory Allocator](#)



## Exploiting Android's Hardened Memory Allocator

Philipp Mao   Elias Valentin Boschung   Marcel Busch   Mathias Payer  
*EPFL, Lausanne, Switzerland*

### Abstract

Most memory corruptions occur on the heap. To harden userspace applications and prevent heap-based exploitation, Google has developed Scudo. Since Android 11, Scudo has replaced jemalloc as the default heap implementation for all native code on Android. Scudo mitigates exploitation attempts of common heap vulnerabilities.

tacks, Scudo implements security measures to ensure the integrity of inline heap metadata and to prevent a predictable heap layout.

Exploitation techniques that target the allocator to escalate a heap-bound memory corruption vulnerability into an arbitrary memory write primitive or code execution have a long tradition. The security community has compiled a large

# Forged Commitbase & Safe Unlink: Zygote to the rescue!

## Both techniques share heavy prerequisites

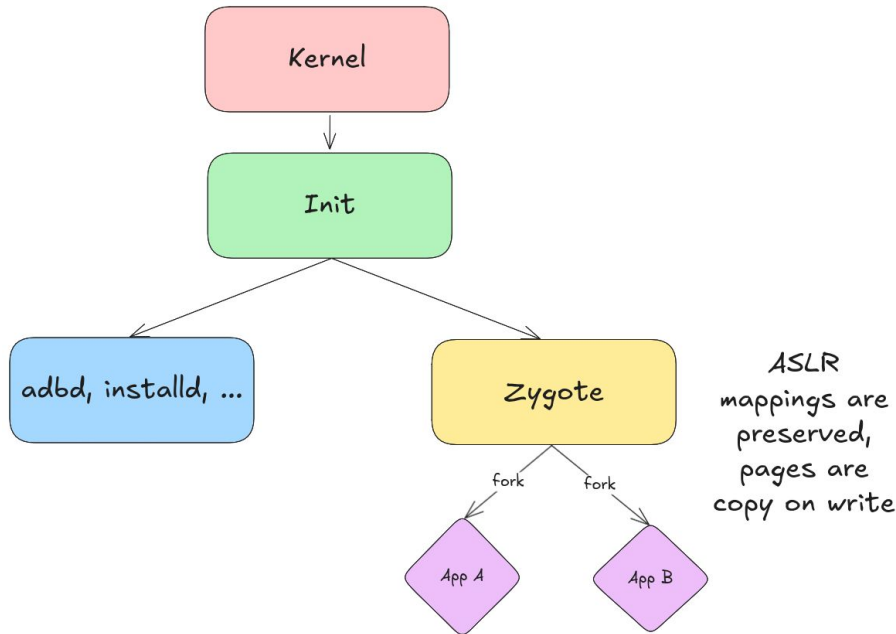
- Ability to guess chunks' addresses
  - Protected by the shuffling of chunks in TransferBatches
- Ability to forge a checksum
  - Protected by the security cookie stored in libc

**But do these protections always work?**

# Forged Commitbase & Safe Unlink: Zygote to the rescue!

## What is Zygote?

- Used to speed up app start time
- Loads framework code/resources
- Forked to spawn app processes
  - All user-installed apps
  - Some privileged services

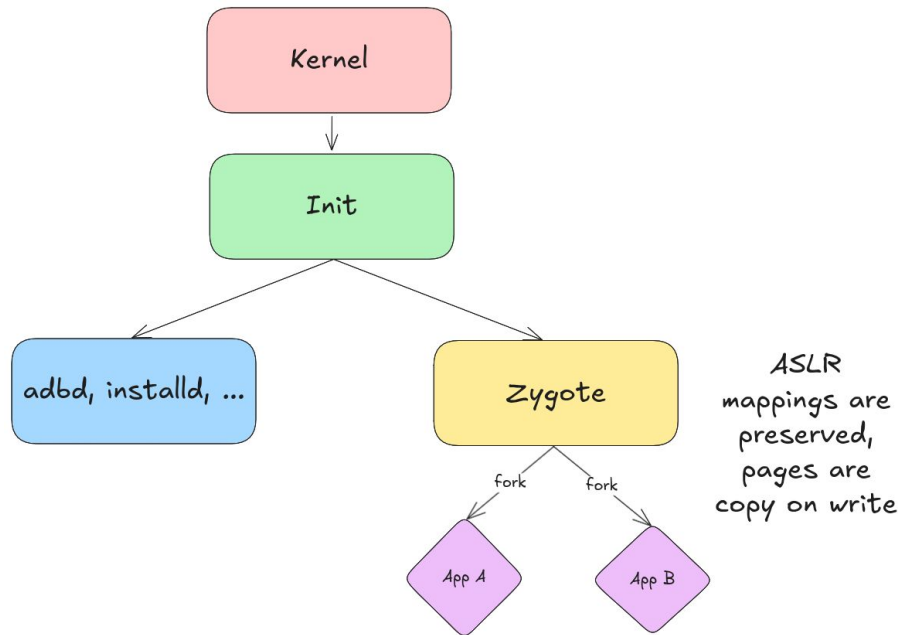


# Forged Commitbase & Safe Unlink: Zygote to the rescue!

## Great! What's the problem then?

- Zygote forks share their memory layout...
- Memory pages are copy on write...
- Zygote uses the heap

➡ Scudo is initialized by Zygote



# Forged Commitbase & Safe Unlink: Zygote to the rescue!

## Both techniques share heavy prerequisites

- Ability to guess chunks' addresses
  - Protected by the shuffling of chunks in TransferBatches

 Random seed is initialized by Zygote!

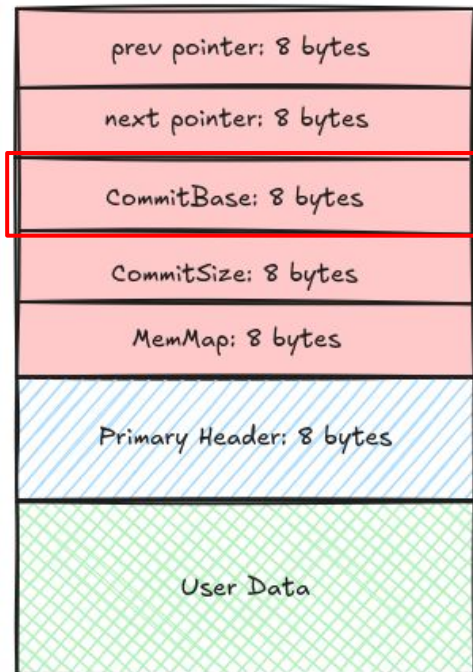
- Ability to forge a checksum
  - Protected by the security cookie stored in libc

 Security cookie is initialized by Zygote!

# Forged CommitBase

Force the **Secondary** allocator to return **one** arbitrary chunk

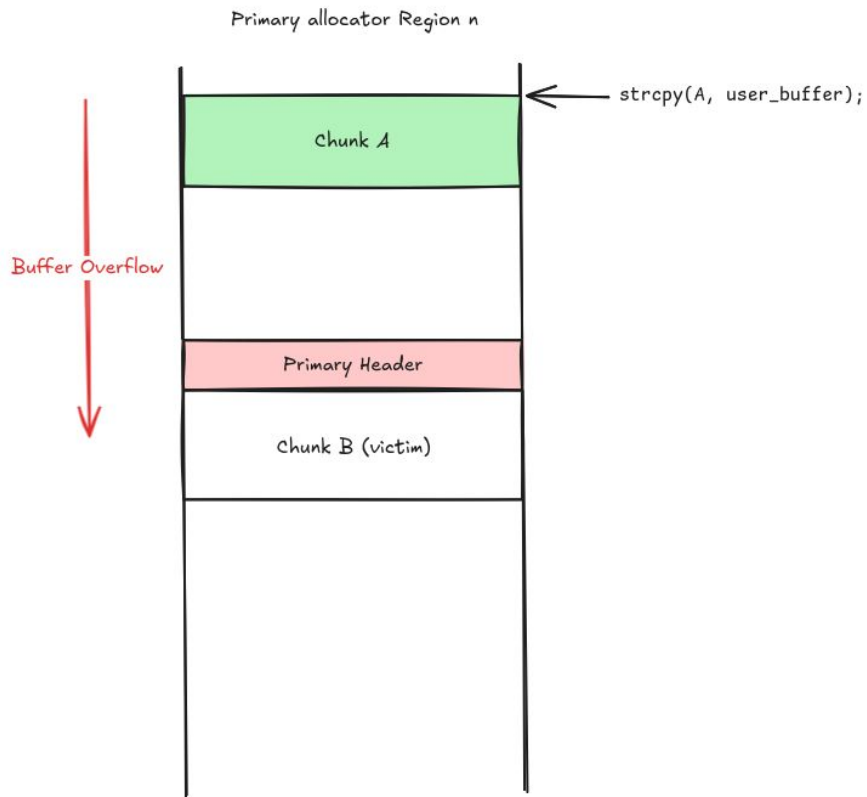
- By forging a fake secondary header
- A few constraints:
  - Predict a chunk's address
  - Forge a checksum
  - Write 0x28 bytes before a chunk
  - Free the victim chunk
  - Have a secondary chunk already allocated





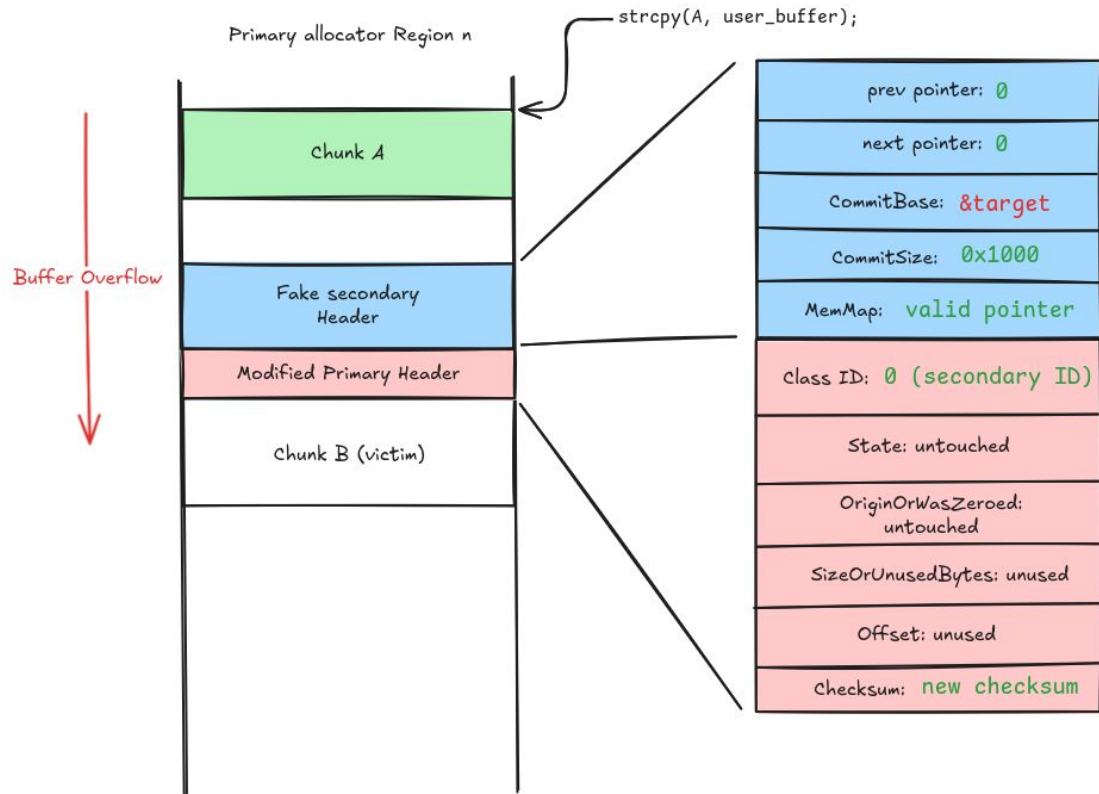
# Forged CommitBase (cont.)

1. Buffer overflow from Chunk A to Chunk B



# Forged CommitBase (cont.)

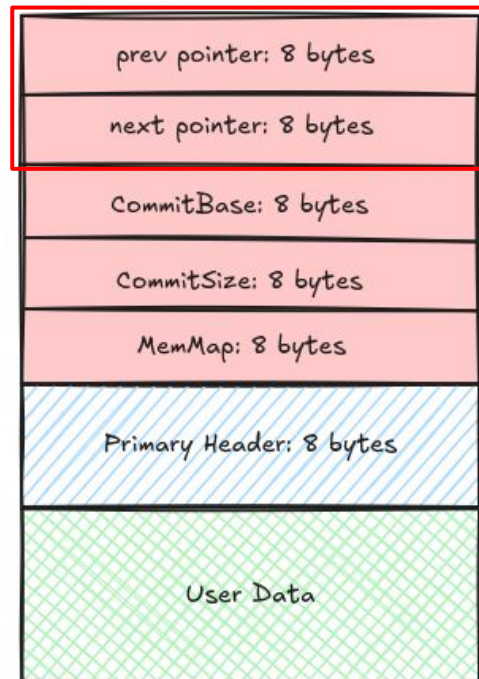
1. Buffer overflow
2. Forge fake secondary header
3. Modify Primary header
4. Free Chunk B
5. **&target** is placed in secondary cache
6. Allocate new secondary chunk & profit



# Safe Unlink

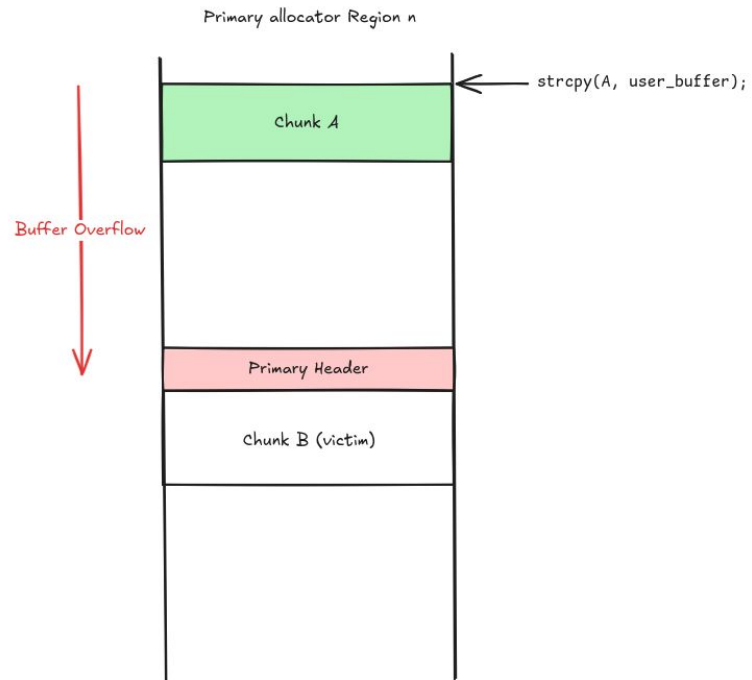

## Force the allocator to return a chunk overlapping the Local Cache

- Inspired by the [Unsafe Unlink](#) technique on GLIBC
- Unlink from the list of in-use chunks
- Once exploited, allows for **unlimited** arbitrary chunks
- Has a few more constraints:
  - Predict a chunk's address
  - Forge a checksum
  - Write 0x30 bytes before a chunk (buffer overflow)
  - Controlled free primitive that we can call 3 times
  - Free the victim chunk



# Safe Unlink (cont.)

*We start with a buffer overflow*




# Safe Unlink (cont.)

1

## Forge fake primary header

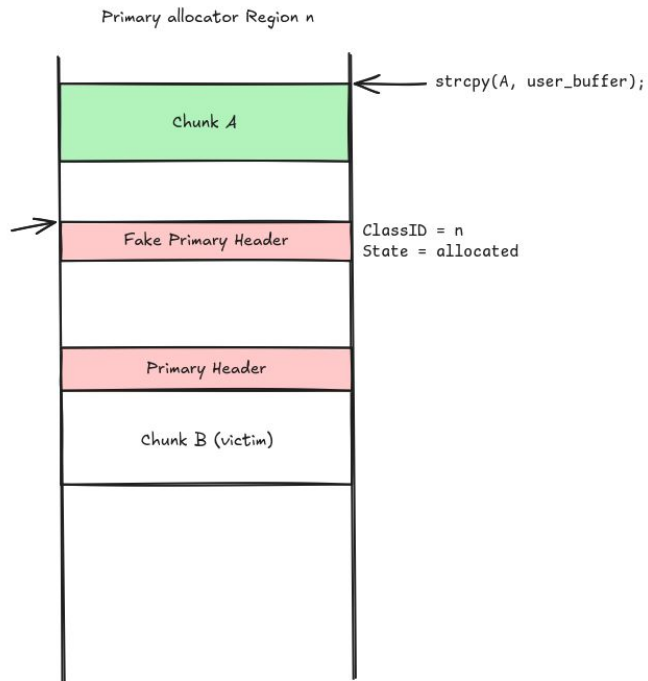
0x30 before victim chunk

PerClass for Region n



1


fake\_header:  
&B - 0x30 bytes



# Safe Unlink (cont.)

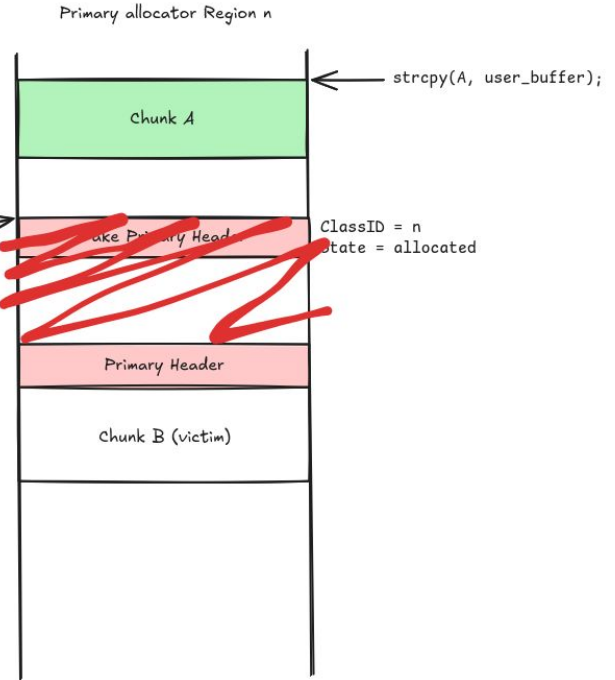
2

**Free the fake chunk**  
Controlled free primitive



2

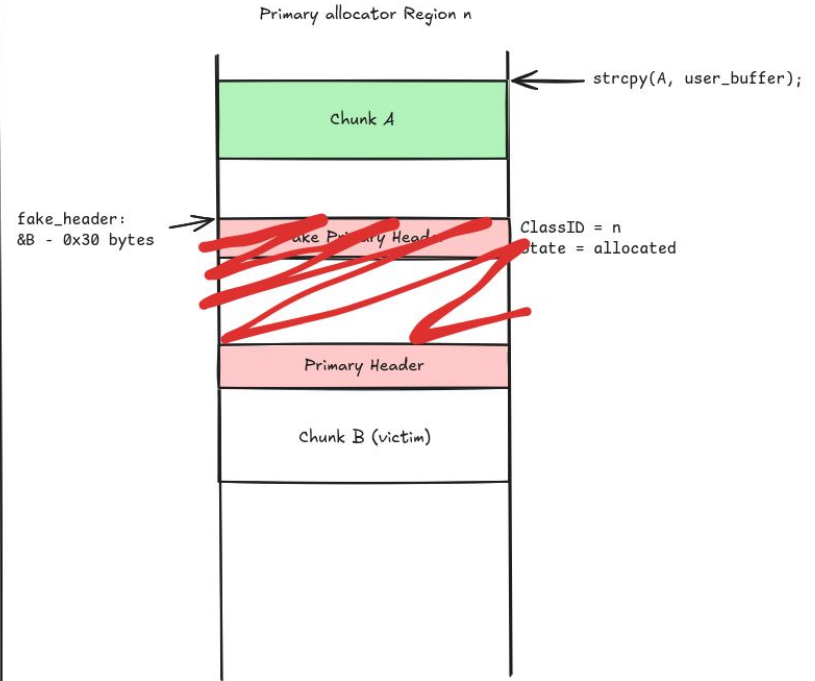

fake\_header:  
&B - 0x30 bytes



# Safe Unlink (cont.)

3

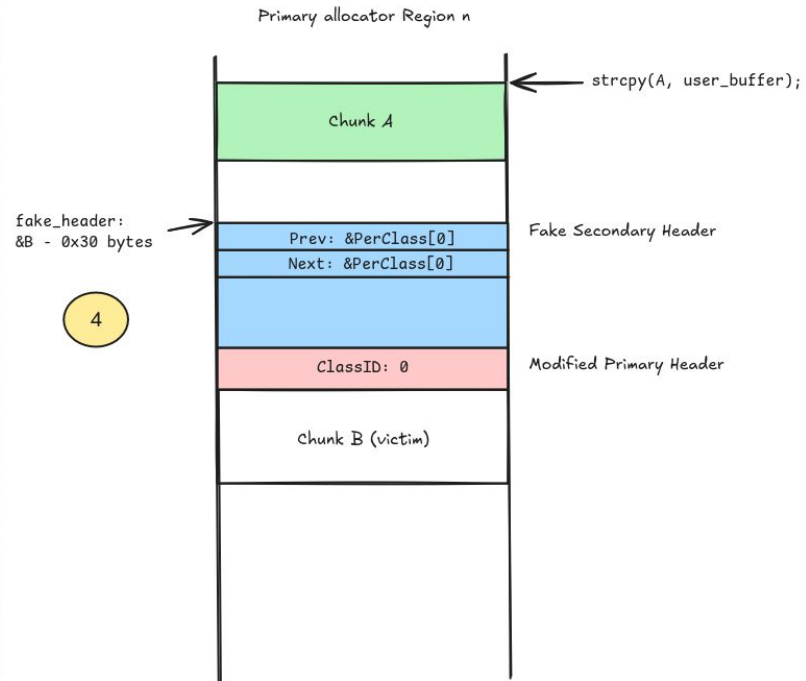

Do this twice



# Safe Unlink (cont.)

4

**Forge fake secondary header**  
Setting up the unlinking attack

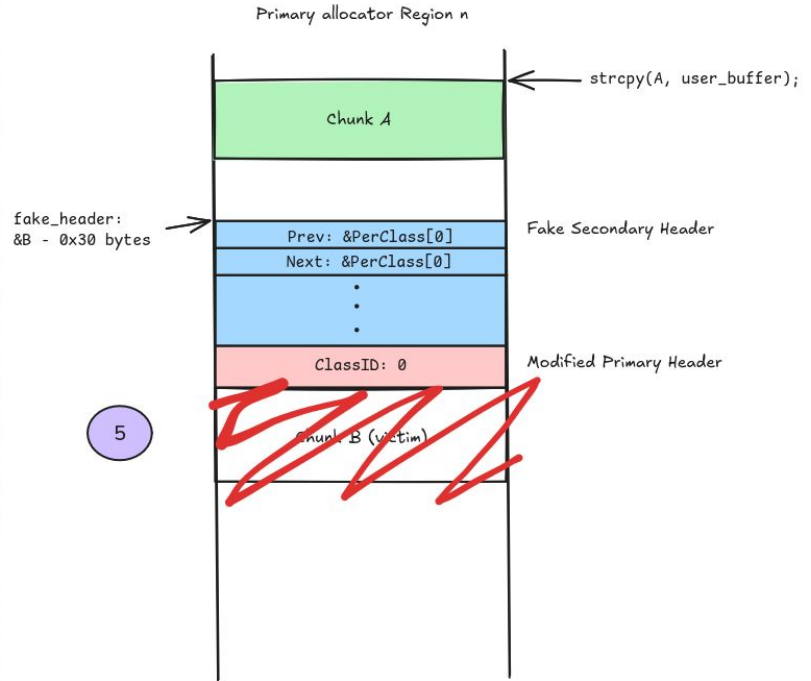





# Safe Unlink (cont.)

5

**Free the victim chunk**  
Trigger the unlinking



# Safe Unlink (cont.)

PerClass for Region n

&fake_header
&fake_header

Fake Secondary Header

Prev: &PerClass[0]
Next: &PerClass[0]
•
•
•

```
typedef struct Node {
    struct Node* prev;
    struct Node* next;
} node_t;

void unlink(node_t* x) {
    node_t* prev = x->prev;
    node_t* next = x->next;

    prev->next = next;
    next->prev = prev;
}
```

# Safe Unlink (cont.)

## Profit!

6

- The next allocation will return a pointer to PerClass
- We now have control over all subsequent allocations

6

PerClass for Region n

&PerClass[0]
&PerClass[0]

# Patches

- Safe Unlink was patched with Compact Pointers in the Local Cache
- CompactPtrs are offsets from the beginning of a Region
- This breaks step 5 of the exploit

```
struct alignas(SCUDO_CACHE_LINE_SIZE) PerClass {  
    u16 Count;  
    u16 MaxCount;  
    // Note: ClassSize is zero for the transfer batch.  
    uptr ClassSize;  
    CompactPtrT Chunks[2 * SizeClassMap::MaxNumCachedHint];  
};
```

# Perspective - Combining the techniques?

*What if we skip the unlinking step?*

Assume we can store arbitrary CompactPtrs in the local cache

- Maybe with Forged CommitBase
- Or another exploit

```
static uptr decompactPtrInternal(uptr Base, CompactPtrT CompactPtr) {  
    return Base + (static_cast<uptr>)(CompactPtr) << CompactPtrScale);  
}
```

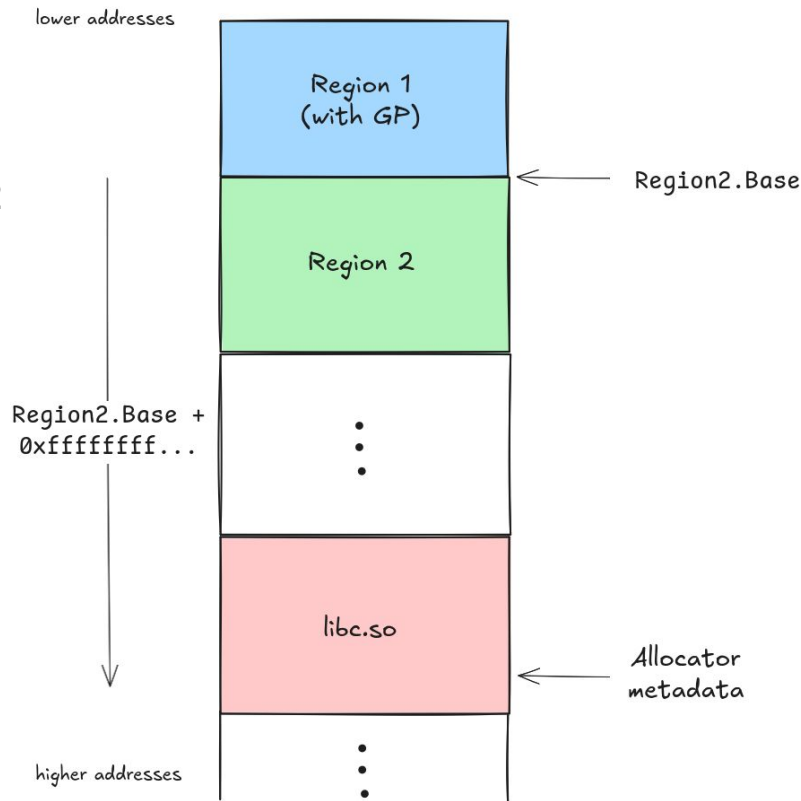
**If we can control this too, we get the same primitive as Safe Unlink!**

# Perspective - Combining the techniques?

*What if we skip the unlinking step?*

We have control over the Local Cache for Region 2

1. Compute distance between Region2 and RegionInfo metadata
2. Write it in Local Cache
3. Allocate chunk over RegionInfo
4. We now control Region2.Base
5. Profit from arbitrary allocations



# Conclusion

- **On Android, Zygote-forked processes can attack each other easily**
  - Forged Commitbase is still unpatched!
  - Safe Unlink has been patched :(
- Combining the techniques can yield interesting results
- In other contexts, Scudo is still very secure!
- There is not much left to exploit (very little inline metadata)
- ARM MTE will mitigate **all** of the exploits we've talked about
  - And many other classes of vulnerabilities
  - It is especially efficient for heap-based memory corruption

# Thank you!

[contact@quarkslab.com](mailto:contact@quarkslab.com)

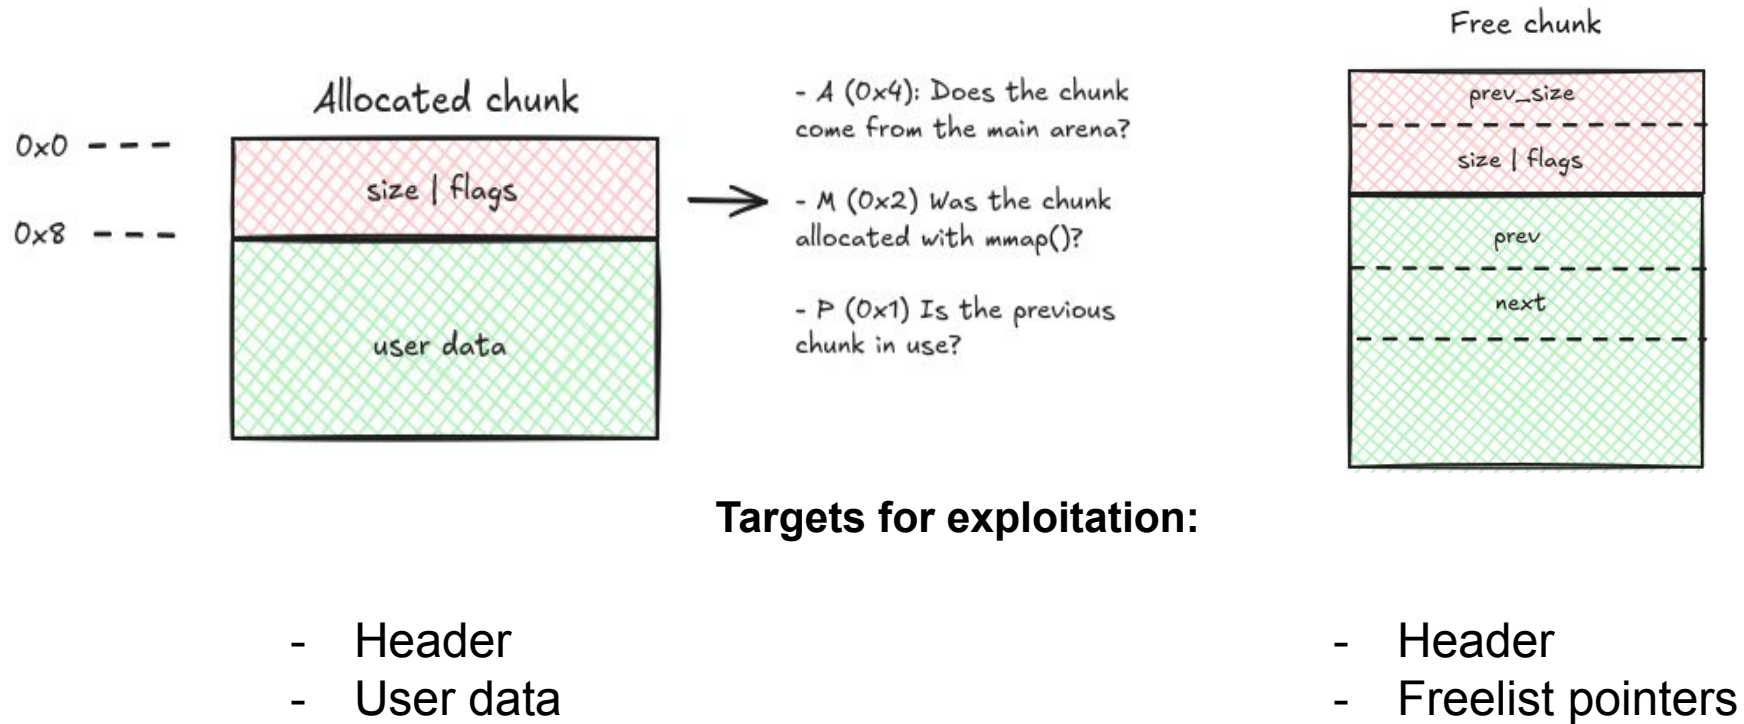
Quarkslab



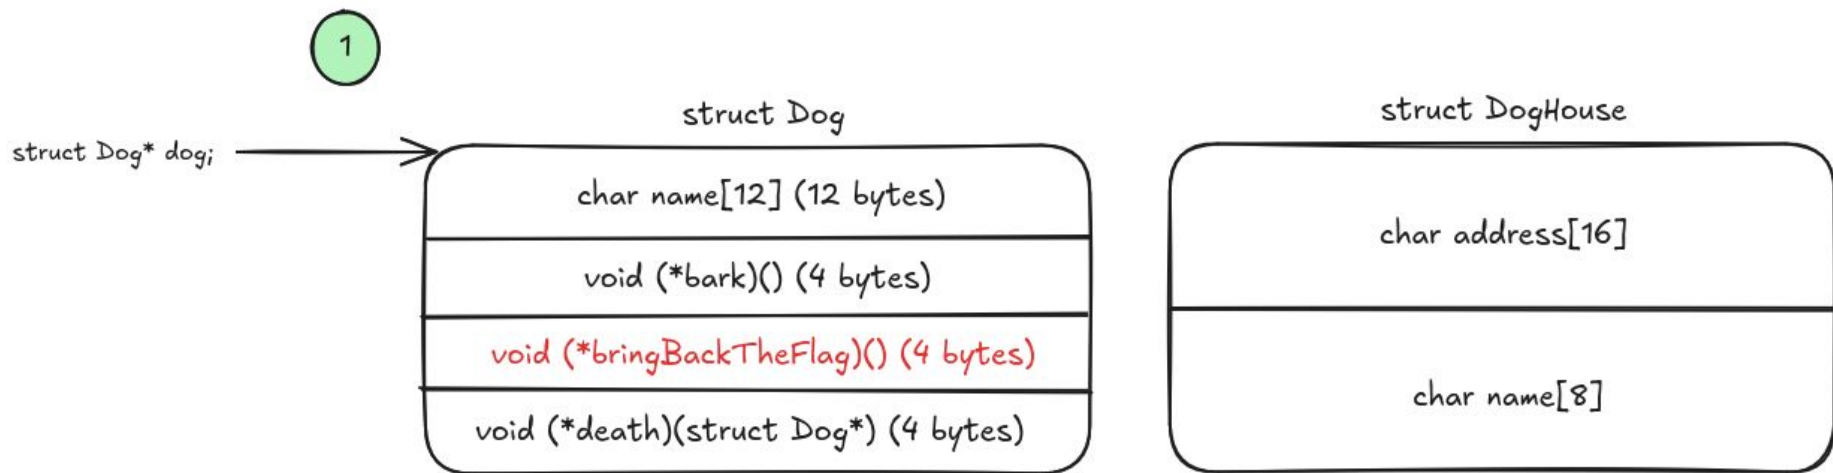
[@philipp0x90](https://twitter.com/philipp0x90)




# GLIBC's ptmalloc2 - what metadata do chunks hold?




# Use-after-free type confusion



# Use-after-free type confusion



# Use-after-free type confusion



# Use-after-free type confusion

